



Writing Optimized and Secure Code

ITP 435 – Spring 2016

Week 5, Lecture 2

Lecturer: Sanjay Madhav

Optimizing Code “As You Go”

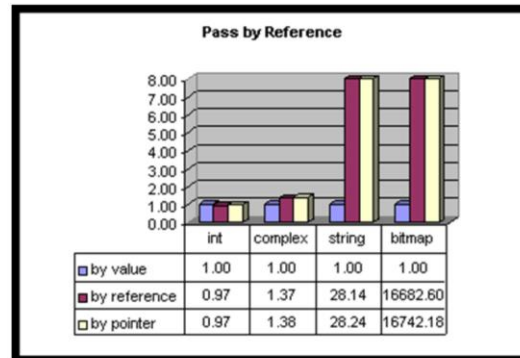


- If you follow some simple rules, you can avoid some common pitfalls.
- These rules will help you write code that's more efficient!

Good Rules to Follow



- Non-basic types should be passed by reference when possible:



From: <http://www.tantalon.com/pete/cppopt/main.htm>

Good Rules to Follow (cont' d)



- Postpone variable declarations when possible:

// Declare Outside (b is true half the time)

```
T x;
```

```
if (b) {
```

```
    x = t;
```

```
    // Do stuff...
```

```
}
```

// Declare Inside (b is true half the time)

```
if (b) {
```

```
    T x = t;
```

```
    // Do stuff...
```

```
}
```

Good Rules to Follow (cont' d)



- Prefer operator+= to operator+ (and like operators)

```
struct Vector2
{
    float X;
    float Y;

    Vector2& operator+=(const Vector2& rhs)
    {
        X += rhs.X;
        Y += rhs.Y;

        return *this;
    }

    Vector2 operator+(const Vector2& rhs)
    {
        Vector2 temp(*this);
        temp += rhs;
        return temp;
    }
};
```

Operator+ has to return a copy by value

Good Rules to Follow (cont' d)



- Use prefix instead of postfix for non-basic types:

```
const T T::operator++ (int) // postfix
{
    T orig(*this); // Have to make copy in postfix
    ++(*this); // call prefix operator
    return (orig);
}
```

Good Rules to Follow (cont' d)



- Prefer explicit constructors

```
class pair
{
    double x, y;
public:
    pair(const string& s) { . . . }
    bool operator == (const pair& c) const { . . . }
};

// Without explicit can do:
pair p;
string s;
if (p == s) { . . . }
```

Good Rules to Follow (cont' d)



- If we change the constructor to explicit

```
explicit pair(const string& s) { . . . }
```

```
// Won't work!
```

```
pair p;
```

```
string s;
```

```
if (p == s) { . . . }
```


80/20 Rule




- “80-20” rule: 80% of your execution time is spent in 20% of your code



Thread-Based Profiling



- If our threads are synchronized we can see if one thread is taking more time than another:



In this case, we clearly want to focus on optimizing our rendering thread first, as it's slowing everything else down.

Profiler




- Profiling programs can allow us to see which functions take up the most time...
- From Intel VTune:

/Function /Call Stack	CPU Time	☆
@ initialize_2D_buffer	11.768s	
@ grid_intersect	5.916s	
@ intersect_objects	5.431s	
@ grid_intersect ← intersect_objects	0.485s	
@ sphere_intersect	5.044s	

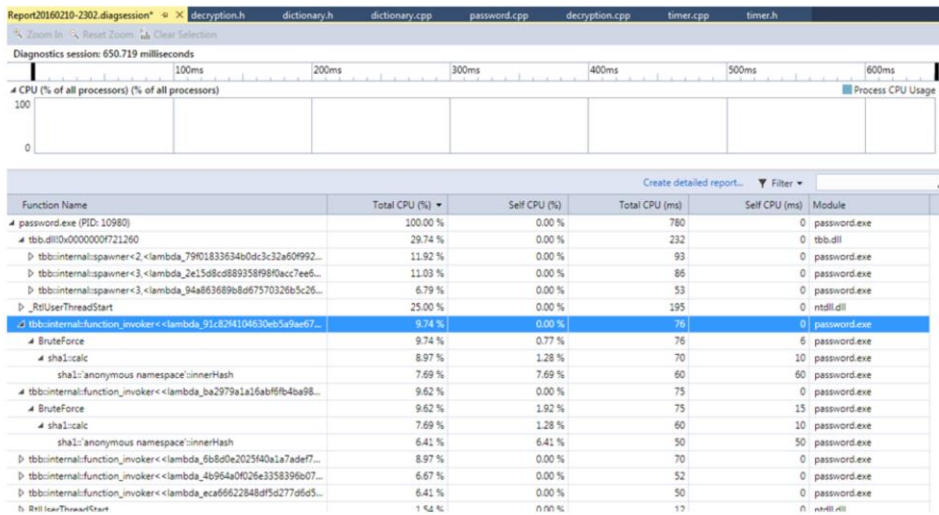
Line	Source	CPU Time	☆
579	cur = g->cells[voxindex];	0.204s	
580	while (cur != NULL) {	0.048s	
581	if (ry->mbox[cur->obj->id] != 2)	1.611s	
582	ry->mbox[cur->obj->id] = ry->	1.025s	
583	cur->obj->methods->intersect(1.098s	



- Product>Profile, select “Time Profiler”



Visual Studio 2015 Profiler



Debug>Start Diagnostic Tools Without Debugging...

How to Optimize?



- Often just algorithmic:
If we're using an $O(n^2)$ solution, maybe there's an $O(n \log n)$ one?
- Also, there are certain rules you can follow while writing your code that will help it be efficient without focused optimizations
- But if your algorithm is fine and you just need to optimize it some more, there are lower level optimizations to consider!

We've already talked about some of our function inlining techniques

Inline Functions



- Tell compiler to copy/paste function code at every location it's called, instead of incurring function call costs
- Method 1: Implementing a function within class declaration suggests you want it inlined:

```
class MyVector3
{
    public:
    float DotProduct(const MyVector3& rhs)
    {
        return (x * rhs.x + y * rhs.y + z * rhs.z);
    }
};
```

You may want to inline a function if it's small and is being called a lot. In this case, there can be a significant performance gain by not calling the function.

Note that virtual functions called virtually *can never be inlined*. This is another hidden cost of virtual functions.

These constructs only *suggest* inlining. The compiler will still ultimately decide whether or not it wants to do it.



- Method 2: For standalone functions, use the “inline” for a suggestion:

```
inline void MyInlineFunction()  
{  
    // Do something  
}
```


Force Inlining



- Visual Studio has a way to “force” inlining, which will inline in almost all cases:

```
__forceinline void MyInlineFunction()  
{  
    // Do something  
}
```

- However, there are still some cases where it won't be inlined:
 - Debug build
 - Recursive function
 - Virtual function called virtually
 - Function pointer call
 - And a couple of other really rare cases

Loop Unrolling



- Regular loop:

```
// This calls an arbitrary function 100 times
for (int x = 0; x < 100; x++)
{
    func();
}
```

If this needs to be high-performance code, the loop as written spends quite a few instructions just incrementing and checking the loop. Also does a lot of branching.

Loop Unrolling, cont' d



- This is a partially “unrolled” version of the same loop:

```
// This calls an arbitrary function 100 times
for (int x = 0; x < 100; x+=5)
{
    func();
    func();
    func();
    func();
    func();
}
```

Disadvantage: Increases size of code

Return Value Optimization



- When returning a class by value, on some compilers it may be more efficient to return an *unnamed* class as opposed to a named one

```
template <class T> T Original(const T& tValue)
{
    T tResult; // named object; optimization potential low
    tResult = tValue;
    return (tResult);
}

template <class T> T Optimized(const T& tValue)
{
    return (T(tValue)); // unnamed; optimization potential high
}
```

Named Return Value Optimization



Many newer versions of compilers support it. From Visual Studio documentation:

“NRVO eliminates the copy constructor and destructor of a stack-based return value. This optimizes out the redundant copy constructor and destructor calls and thus improves overall performance. It is to be noted that **this could lead to different behavior between optimized and non-optimized programs.**”

Named Return Value Optimization (cont' d)



```
#include <stdio.h>
struct RVO
{
    RVO(){printf("I am in constructor\n");}
    RVO (const RVO& c_RVO) {printf ("I am in copy constructor\n");}
    ~RVO(){printf ("I am in destructor\n");}
    int mem_var;
};
RVO MyMethod (int i)
{
    RVO rvo;
    rvo.mem_var = i;
    return (rvo);
}
int main()
{
    RVO rvo;
    rvo=MyMethod(5);
}
```

NRVO Output



- Without NRVO (cl /Od sample1.cpp), the expected output would be:

```
I am in constructor  
I am in constructor  
I am in copy constructor  
I am in destructor  
I am in destructor  
I am in destructor
```

- With NRVO (cl /O2 sample1.cpp), the expected output would be:

```
I am in constructor  
I am in constructor  
I am in destructor  
I am in destructor
```

This code works but...



```
#include <iostream>
int main() {
    bool allowAccess = false;
    char buffer[16];
    std::cout << "Enter Password:";
    std::cin >> buffer;

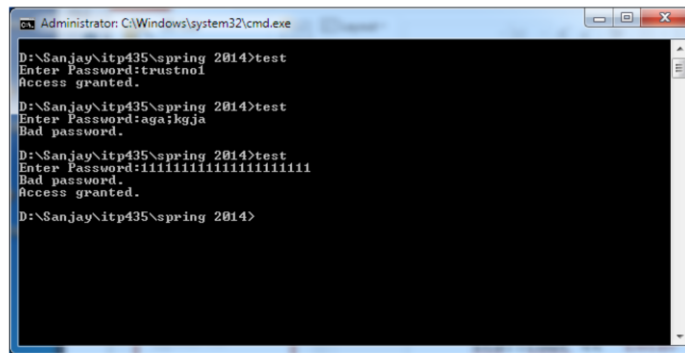
    if (std::strcmp(buffer, "trustno1") == 0) {
        allowAccess = true;
    } else {
        std::cout << "Bad password." << std::endl;
    }

    if (allowAccess) {
        std::cout << "Access granted." << std::endl;
        // DO STUFF
    }

    return 0;
}
```

Time for a demonstration...

Some sample input to the program...



```
Administrator: C:\Windows\system32\cmd.exe
D:\Sanjay\itp435\spring 2014>test
Enter Password:trustno1
Access granted.
D:\Sanjay\itp435\spring 2014>test
Enter Password:aga;kgja
Bad password.
D:\Sanjay\itp435\spring 2014>test
Enter Password:iiiiiiiiiiiiiiiiiiii
Bad password.
Access granted.
D:\Sanjay\itp435\spring 2014>
```

Why? Well, first let me show you something...



26

What if we look at the code in a disassembler?



NUJ

; Attributes: bp-based frame

sub_1285370 proc near

var_14= byte ptr -14h

var_1= byte ptr -1

```
push    ebp
mov     ebp, esp
sub     esp, 54h
push    ebx
push    esi
push    edi
mov     [ebp+var_1], 0
push    offset aEnterPassword ; "
mov     eax, ds:?cout@std@@@3U?$.
push    eax
call    sub_1281302
add     esp, 8
lea     eax, [ebp+var_14]
push    eax
mov     ecx, ds:?cin@std@@@3U?$.basic_istream
push    ecx
call    sub_128100C
add     esp, 8
push    offset aTrustno1 ; "trustno1"
lea     eax, [ebp+var_14]
push    eax
call    j_strcmp
add     esp, 8
test    eax, eax
jnz     short loc_128538E
```

Hmm...So there are 2
variables on the stack,
separated by 0x13
bytes

USC Viterbi
School of Engineering

University of Southern California

I don't expect you to know the assembly, of course

What if we look at the code in a disassembler?



GNU

; Attributes: bp-based frame

sub_1285370 proc near

var_14= byte ptr -14h

var_1= byte ptr -1

push ebp

mov ebp, esp

sub esp, 54h

push ebx

push esi

push edi

mov [ebp+var_1], 0

push offset aEnterPassw ; "Enter Password" ; std::basic_ostream<char,std::

eax, ds:7c00tstd@00401000 ; std::basic_ostream<char,std::

push eax

call sub_1281302

add esp, 8

lea eax, [ebp+var_14]

push eax

mov ecx, ds:7c00cinstd@00401000 ; std::basic_istream<char,std::

push ecx

call sub_128100C

add esp, 8

push offset aTrustno1 ; "trustno1"

lea eax, [ebp+var_14]

push eax ; char *

call j_strcmp

add esp, 8

test eax, eax

jnz short loc_12853BE

var_14 is read
from a cin...

I don't expect you to know the assembly, of course

What if we look at the code in a disassembler?



```
sub_1285370 proc near
var_14= byte ptr -14h
var_1= byte ptr -1

push    ebp
mov     ebp, esp
sub     esp, 54h
push    ebx
push    esi
push    edi
mov     [ebp+var_1], 0
push    offset aEnterP
mov     eax, ds:coutd
push    eax
call    sub_1281302
add     esp, 8
lea     eax, [ebp+var_14]
push    eax
mov     ecx, ds:cin@std@0307$basic_istream
push    ecx
call    sub_128100C
add     esp, 8
push    offset aTrustno1 ; "trustno1"
lea     eax, [ebp+var_14]
push    eax
call    j_strcmp
add     esp, 8
test    eax, eax
jnz     short loc_12853BE
```

strcmp is being used,
and I can clearly see
that var_14 is a char *

What if we look at the code in a disassembler?



Access seems
to be granted
based on
var_1

```
offset sub_1281466  
push offset aAccessGranted ; "Access granted."  
push  
mov eax, ds:cout@std::basic_ostream<char, std::char_traits<char>>::cout  
push eax
```

```
loc_12853DE:  
movzx eax, [ebp+var_1]  
test eax, eax  
jz short loc_1285406
```

Stack Variable Allocation



- In our program's case, the stack variables are something like this:

Address	
0x00	char buffer[16]
0x04	
0x08	
0x0C	
0x10	???padding???
0x14	bool allowAccess

Normal usage case...



- If we write 15 or less characters to buffer, we're okay:

Address				
0x00	't'	'r'	'u'	's'
0x04	't'	'n'	'o'	'1'
0x08	'\0'			
0x0C				
0x10	???padding???			
0x14	1	0	0	0

"Malformed string" Usage Case



- If we write exactly 21 letters we'll get something like:

Address				
0x00	'1'	'1'	'1'	'1'
0x04	'1'	'1'	'1'	'1'
0x08	'1'	'1'	'1'	'1'
0x0C	'1'	'1'	'1'	'1'
0x10	'1'	'1'	'1'	'1'
0x14	'1'	'\0'	0	0

The least significant byte of allowAccess is now non-zero, which is true!!!

How to fix this problem?



```
#include <iostream>
#include <string>
int main() {
    bool allowAccess = false;
    std::string buffer;
    std::cout << "Enter Password:";
    std::cin >> buffer;

    if (buffer == "trustno1") {
        allowAccess = true;
    } else {
        std::cout << "Bad password." << std::endl;
    }

    if (allowAccess) {
        std::cout << "Access granted." << std::endl;
        // DO STUFF
    }
    return 0;
}
```

The easy answer is use `std::string` instead, as it cannot be susceptible to stack-based buffer overruns, since the storage is on the heap and it dynamically will resize if too many characters in the input

If you really must use c-style strings...



- In Visual Studio, there are so-called “Secure” CRT functions
[http://msdn.microsoft.com/en-us/library/8ef0s5kh\(v=vs.90\).aspx](http://msdn.microsoft.com/en-us/library/8ef0s5kh(v=vs.90).aspx)

- For example, we could change our previous example to be:

```
void foo (char *bar)
{
    char c[12];
    strcpy_s(c, bar);
}
```

- If it looks the same, the secret is in templates:

```
template <size_t size>
errno_t strcpy_s(
    char (&strDestination)[size],
    const char *strSource );
```

[http://msdn.microsoft.com/en-us/library/8ef0s5kh\(v=vs.90\).aspx](http://msdn.microsoft.com/en-us/library/8ef0s5kh(v=vs.90).aspx)

And buffer overruns can be a lot worse



- A more complex, but possible buffer overrun exploit is to execute arbitrary code injected by the attacker
- **shellcode** can be used in a string to invoke a shell session with (ideally) admin/root access:

```
char shellcode[ ] = "\x31\xc0\x50\x68//sh\x68/bin\x89\xe3\x50\x53\x89\xe1\x99\xb0\x0b\xcd\x80";
```

Double Free Vulnerability



- Freeing/deleting a pointer twice corrupts the underlying heap and can potentially allow an attacker to cause a buffer overflow

```
char* ptr = new char[SIZE];  
...  
if (abrt)  
{  
    delete[] ptr;  
}  
...  
delete[] ptr;
```

Double Free Vulnerability (cont' d)



- Solution: Set pointer to nullptr when deleted. It is safe to delete a nullptr.

```
char* ptr = new char[SIZE];  
...  
if (abrt)  
{  
    delete[] ptr;  
    ptr = nullptr; // VULNERABILITY FIXED  
}  
...  
delete[] ptr;
```

Some Interesting Security Links



- <https://nakedsecurity.sophos.com/2014/02/24/anatomy-of-a-goto-fail-apples-ssl-bug-explained-plus-an-unofficial-patch/>
- <http://arstechnica.com/security/2015/12/researchers-confirm-backdoor-password-in-juniper-firewall-code/>
- <https://code.google.com/p/google-security-research/issues/detail?id=675>
- <http://arstechnica.com/security/2015/08/how-security-flaws-work-the-buffer-overflow/>
- <http://thehackernews.com/2014/04/heartbleed-bug-explained-10-most.html>

<https://nakedsecurity.sophos.com/2014/02/24/anatomy-of-a-goto-fail-apples-ssl-bug-explained-plus-an-unofficial-patch/>

<http://arstechnica.com/security/2015/12/researchers-confirm-backdoor-password-in-juniper-firewall-code/>

<https://code.google.com/p/google-security-research/issues/detail?id=675>

<http://arstechnica.com/security/2015/08/how-security-flaws-work-the-buffer-overflow/>

<http://thehackernews.com/2014/04/heartbleed-bug-explained-10-most.html>